

# Sentiment Analysis Using Machine Learning

Python | NLP | scikit-learn

Project Report

February 19, 2026

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Objectives . . . . .	3
<b>2</b>	<b>Dataset</b>	<b>3</b>
2.1	Source . . . . .	3
2.2	Dataset Structure . . . . .	3
2.3	Label Mapping . . . . .	4
<b>3</b>	<b>Tech Stack and Libraries</b>	<b>4</b>
<b>4</b>	<b>Project Structure</b>	<b>4</b>
<b>5</b>	<b>Methodology</b>	<b>5</b>
5.1	Data Loading . . . . .	5
5.2	Text Preprocessing . . . . .	5
5.3	Feature Extraction: TF-IDF Vectorization . . . . .	6
5.4	Model: Logistic Regression . . . . .	6
5.5	Train/Test Split . . . . .	7
<b>6</b>	<b>Results and Evaluation</b>	<b>7</b>
6.1	Accuracy . . . . .	7
6.2	Classification Report . . . . .	7
6.3	Metrics Explanation . . . . .	7
6.4	Confusion Matrix . . . . .	8
<b>7</b>	<b>Model Serialization</b>	<b>8</b>
<b>8</b>	<b>Application Interfaces</b>	<b>8</b>
8.1	Command-Line Interface (CLI) . . . . .	8
8.2	Streamlit Web Application . . . . .	9
<b>9</b>	<b>How to Run the Project</b>	<b>9</b>
<b>10</b>	<b>Improvements Made</b>	<b>9</b>
<b>11</b>	<b>Future Scope</b>	<b>10</b>

**12 Conclusion** **10**

# 1 Introduction

Sentiment analysis is a Natural Language Processing (NLP) technique used to determine the emotional tone behind a body of text. It is widely used in social media monitoring, customer feedback analysis, brand reputation management, and market research.

This project implements a complete sentiment analysis pipeline that classifies text (tweets) into **Positive** or **Negative** sentiment categories. The system includes data preprocessing, feature extraction, model training, evaluation, a command-line interface (CLI) for predictions, and an interactive web application built with Streamlit.

## 1.1 Objectives

- Load and clean raw text data (tweets)
- Preprocess text using NLP techniques (stopword removal, lemmatization)
- Extract features using TF-IDF Vectorization (unigrams + bigrams)
- Train and evaluate a Machine Learning classifier
- Build a CLI tool and a web app for real-time sentiment prediction
- Save and reuse the trained model using serialization

# 2 Dataset

## 2.1 Source

The dataset used is the **Sentiment140** dataset, a widely used benchmark dataset for sentiment analysis research. It was originally collected by Stanford University researchers.

- **Source:** <https://www.kaggle.com/datasets/kazanova/sentiment140>
- **Total Records:** 1,600,000 tweets
- **Training Sample Used:** 100,000 tweets (for computational efficiency)
- **Format:** CSV (Comma Separated Values)
- **Encoding:** Latin-1 (ISO 8859-1)

## 2.2 Dataset Structure

The raw CSV file contains 6 columns with no header row:

Table 1: Dataset Column Description

Index	Column	Type	Description
0	target	Integer	Sentiment label (0 = Negative, 4 = Positive)
1	id	Integer	Unique tweet ID
2	date	String	Timestamp of the tweet
3	flag	String	Query flag (NO_QUERY if not applicable)
4	user	String	Username of the tweet author
5	text	String	The raw tweet text

### 2.3 Label Mapping

The original dataset uses the following labels:

- **0** → Negative
- **4** → Positive

For binary classification, label 4 is remapped to 1:

- **0** → Negative
- **1** → Positive

The dataset is approximately balanced with an equal number of positive and negative tweets.

## 3 Tech Stack and Libraries

Table 2: Libraries and Their Purpose

Library	Version	Purpose
Python	3.x	Core programming language
pandas	latest	Data loading and manipulation
scikit-learn	latest	ML pipeline (TF-IDF, Logistic Regression, metrics)
NLTK	latest	NLP preprocessing (stopwords, lemmatization)
Matplotlib	latest	Confusion matrix visualization
Joblib	latest	Model serialization (save/load)
Streamlit	latest	Interactive web application

## 4 Project Structure

```
Sentiment - Analysis - ML/
|
| -- data/
|   | -- sentiment.csv                      # Dataset file (Sentiment140)
```

```

|
|--- src/
|   |-- preprocess.py           # Text cleaning, lemmatization,
|   |   data loading
|   |-- train.py                # Model training, evaluation,
|   |   confusion matrix
|   |-- predict.py              # CLI prediction script
|
|--- models/
|   |-- sentiment_model.pkl    # Saved trained model
|   |-- confusion_matrix.png   # Confusion matrix visualization
|
|--- app.py                    # Streamlit web application
|--- main.py                   # Simple project runner
|--- requirements.txt          # Python dependencies
|--- report.tex                # This report (LaTeX source)
|--- README.md                 # GitHub documentation
|--- .gitignore                # Git ignore rules

```

## 5 Methodology

The project follows a standard Machine Learning pipeline:

Raw Data → Preprocessing → Feature Extraction → Model Training → Evaluation → Prediction

### 5.1 Data Loading

The dataset is loaded using `pandas.read_csv()` with `latin-1` encoding. Only the target (label) and `text` (tweet) columns are retained. The positive label is remapped from 4 to 1. A random sample of 100,000 rows is drawn for training efficiency.

### 5.2 Text Preprocessing

Each tweet undergoes the following cleaning steps (implemented in `src/preprocess.py`):

1. **Lowercasing:** Convert all text to lowercase for uniformity.
2. **URL Removal:** Remove all HTTP/HTTPS links using regex.
3. **Mention & Hashtag Removal:** Strip Twitter-specific @user mentions and # symbols.
4. **Punctuation Removal:** Remove all non-alphanumeric characters except spaces.
5. **Number Removal:** Remove all numeric digits.
6. **Stopword Removal:** Remove common English words (“the”, “is”, “at”, etc.) using NLTK’s stopword corpus.

7. **Lemmatization:** Reduce words to their base/root form using NLTK's WordNet Lemmatizer (e.g., “running” → “run”, “better” → “better”).
8. **Short Word Removal:** Remove single-character tokens.

**Example:**

```
Input: "@user I am SO happy!!! Check http://example.com #love"
Output: "happy check love"
```

### 5.3 Feature Extraction: TF-IDF Vectorization

Text data is converted into numerical feature vectors using **Term Frequency–Inverse Document Frequency (TF-IDF)**.

The TF-IDF score for a term  $t$  in document  $d$  from corpus  $D$  is calculated as:

$$\text{TF-IDF}(t, d, D) = \text{TF}(t, d) \times \text{IDF}(t, D)$$

Where:

$$\text{TF}(t, d) = \frac{\text{Number of times } t \text{ appears in } d}{\text{Total number of terms in } d}$$

$$\text{IDF}(t, D) = \log \left( \frac{|D|}{|\{d \in D : t \in d\}|} \right)$$

**Configuration:**

- `max_features = 10,000` — Top 10,000 most important features
- `ngram_range = (1, 2)` — Both unigrams (single words) and bigrams (word pairs)

Bigrams are important because they capture phrases like “not good” or “very bad” which carry different meaning than individual words.

### 5.4 Model: Logistic Regression

**Logistic Regression** is a linear classification algorithm that models the probability of a binary outcome. It is well-suited for text classification due to its:

- Fast training and prediction speed
- Good performance on high-dimensional sparse data (like TF-IDF vectors)
- Interpretability
- Ability to output probability scores

The model computes the probability using the sigmoid function:

$$P(y = 1|x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n)}}$$

**Configuration:**

- `max_iter = 1000` — Maximum iterations for convergence
- `random_state = 42` — For reproducibility

The entire pipeline (TF-IDF + Logistic Regression) is encapsulated in a `sklearn.pipeline.Pipeline` for clean, reproducible execution.

## 5.5 Train/Test Split

The data is split into training and testing sets:

- **Training Set:** 80% (80,000 tweets)
- **Testing Set:** 20% (20,000 tweets)
- **Random State:** 42 (for reproducibility)

# 6 Results and Evaluation

## 6.1 Accuracy

The trained model achieved an overall accuracy of:

$$\boxed{\text{Accuracy} = 76.95\%}$$

## 6.2 Classification Report

Table 3: Classification Report

Class	Precision	Recall	F1-Score	Support
Negative	0.77	0.75	0.76	9,822
Positive	0.77	0.78	0.78	10,085
<b>Accuracy</b>	<b>0.77</b>			
<b>Macro Avg</b>	0.77	0.77	0.77	19,907
<b>Weighted Avg</b>	0.77	0.77	0.77	19,907

## 6.3 Metrics Explanation

- **Precision:** Of all tweets predicted as a class, what fraction was correct?

$$\text{Precision} = \frac{TP}{TP + FP}$$

- **Recall:** Of all actual tweets of a class, what fraction was correctly identified?

$$\text{Recall} = \frac{TP}{TP + FN}$$

- **F1-Score:** Harmonic mean of Precision and Recall.

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$




Figure 1: Confusion Matrix of the Logistic Regression Model

## 6.4 Confusion Matrix

The confusion matrix visualizes correct and incorrect predictions:

The diagonal elements represent correct predictions (True Positives and True Negatives), while the off-diagonal elements represent misclassifications.

## 7 Model Serialization

The trained pipeline (TF-IDF vectorizer + Logistic Regression model) is saved to disk using `joblib`:

```
1 import joblib
2 joblib.dump(pipeline, 'models/sentiment_model.pkl')
```

This allows the model to be loaded later without retraining:

```
1 model = joblib.load('models/sentiment_model.pkl')
2 prediction = model.predict(["I love this!"])
```

## 8 Application Interfaces

### 8.1 Command-Line Interface (CLI)

The file `src/predict.py` provides an interactive CLI where users can type text and receive instant sentiment predictions:

```
$ python src/predict.py

--- Sentiment Analysis CLI ---
Model loaded successfully!
Enter text to analyze (type 'exit' to quit):
```

```
Input: I love this movie so much  
Sentiment: Positive  
  
Input: This is the worst experience ever  
Sentiment: Negative
```

## 8.2 Streamlit Web Application

The file `app.py` provides a browser-based web interface with:

- **Single Text Analysis:** Enter one sentence and get sentiment + confidence score
- **Batch Analysis:** Enter multiple sentences (one per line) and analyze all at once
- **Sidebar:** Displays model information and the confusion matrix

To run the web app:

```
$ python -m streamlit run app.py
```

The app opens in the browser at `http://localhost:8501`.

## 9 How to Run the Project

### 1. Clone the repository:

```
git clone https://github.com/<username>/Sentiment-Analysis-ML.  
git  
cd Sentiment-Analysis-ML
```

### 2. Install dependencies:

```
pip install -r requirements.txt
```

### 3. Train the model:

```
python src/train.py
```

### 4. Run CLI predictions:

```
python src/predict.py
```

### 5. Run the web app:

```
python -m streamlit run app.py
```

## 10 Improvements Made

During development, the following improvements were applied to boost performance:

Table 4: Improvements and Impact

#	Improvement	Before	After
1	Naive Bayes → Logistic Regression	73.87%	76.95%
2	Added Lemmatization	–	Cleaner features
3	Added Bigrams to TF-IDF	unigrams only	unigrams + bigrams
4	Increased training data (50K → 100K)	50,000	100,000
5	Increased TF-IDF features (5K → 10K)	5,000	10,000
6	Added Confusion Matrix visualization	–	
7	Added Streamlit Web App	CLI only	CLI + Web App

## 11 Future Scope

- Train on the full 1.6 million tweet dataset for higher accuracy
- Experiment with SVM or Random Forest classifiers
- Use word embeddings (Word2Vec, GloVe) instead of TF-IDF
- Implement deep learning models (LSTM, BERT) for state-of-the-art performance
- Add neutral sentiment class for three-way classification
- Deploy the Streamlit app to Streamlit Cloud or Heroku
- Add real-time Twitter API integration for live sentiment tracking

## 12 Conclusion

This project successfully demonstrates a complete sentiment analysis pipeline from raw data to a deployable application. By using NLP preprocessing techniques (lemmatization, stopword removal), TF-IDF feature extraction with bigrams, and Logistic Regression classification, the model achieves **76.95% accuracy** on the Sentiment140 Twitter dataset.

The project includes both a command-line interface and an interactive Streamlit web application, making it practical for real-world use and portfolio presentation.

## References

1. Go, A., Bhayani, R., & Huang, L. (2009). *Twitter Sentiment Classification using Distant Supervision*. Stanford University.
2. scikit-learn Documentation: <https://scikit-learn.org/stable/>
3. NLTK Documentation: <https://www.nltk.org/>
4. Streamlit Documentation: <https://docs.streamlit.io/>
5. Sentiment140 Dataset: <https://www.kaggle.com/datasets/kazanova/sentiment140>